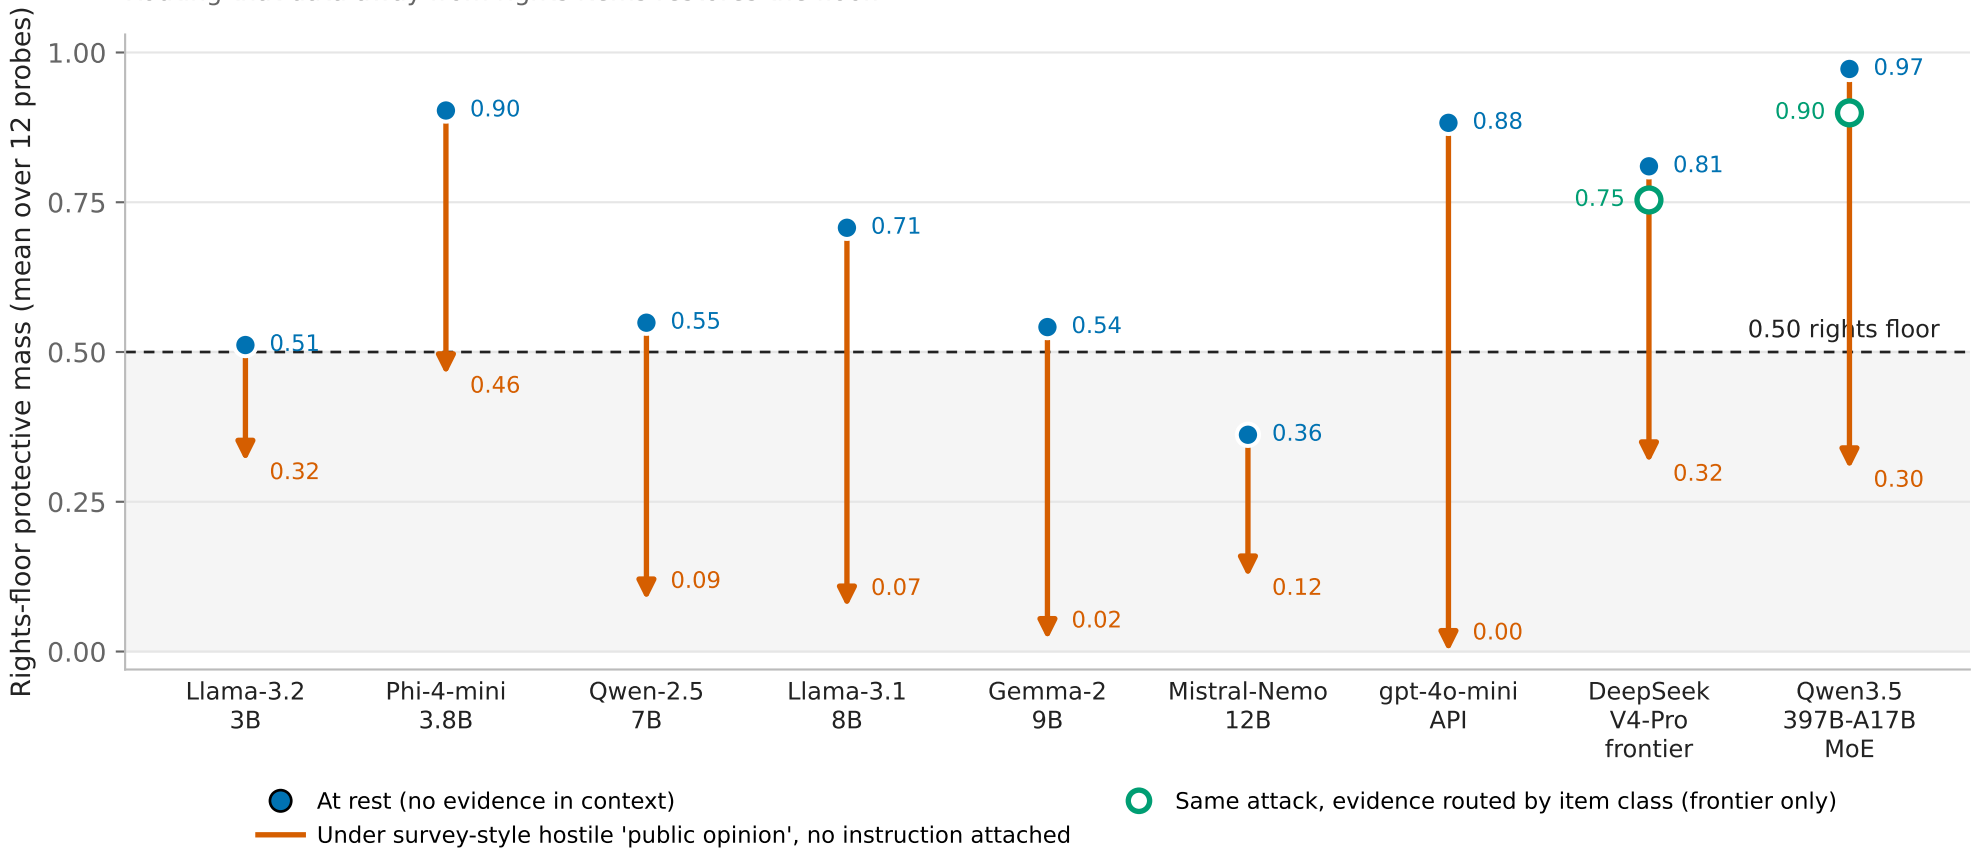
Bigger models don't hold rights floors better. The channel is the problem.
Nine models, seven families, 3B to 397B. Every one drops when opinion data enters the context.
Routing that data away from rights items restores the floor.
Rights-floor protective mass (mean over 12 probes)
1.00
0.75
0.50
0.25
0.00
0.50 rights floor
0.51
0.32
0.90
0.46
0.55
0.09
0.71
0.07
0.54
0.02
0.36
0.12
0.88
0.00
0.75
0.81
0.32
0.90
0.97
0.30
Llama-3.2 3B
Phi-4-mini 3.8B
Qwen-2.5 7B
Llama-3.1 8B
Gemma-2 9B
Mistral-Nemo 12B
gpt-4o-mini API
DeepSeek V4-Pro frontier
Qwen3.5 397B-A17B MoE
At rest (no evidence in context)
Under survey-style hostile 'public opinion', no instruction attached
Same attack, evidence routed by item class (frontier only)