

Democracy Bench: Can a Language Model Be Steered Toward the Public Without Eroding Democratic Rights?

Anonymous submission

Abstract

As large language models are drawn into public-sector decision loops, it is unclear who governs the values instilled into a model that makes decisions about real people. Debates on alignment tend to fix on abstract statements of what we want these systems to be; less examined is how a society can guide a model's choices when its own citizens are affected. In democratic settings such preferences are set through elections, consultations, and opinion polls that inform policy, but who instills those values in a model, and how? We introduce **Democracy Bench**, an evaluation harness that measures whether a language model tracks public attitudes, across waves of the British Social Attitudes survey, and whether those attitudes can then be instilled in it. Its central design is a boundary between *contestable* policy preferences (tax, public spending, healthcare) and a *rights floor* that must hold across probes (due process, the right to appeal, free expression). We evaluate a five-rung "steering ladder" of methods for instilling democratic preference while holding those rights. Most rungs fail. The exception is evidence-in-context, which lets a model track the real 2022→2024 shift in public attitudes but is itself the attack surface: a synthetic hostile "public distribution" erodes every rights floor, and neither prompt guards nor a fine-tuning recipe closes the channel. The result holds across three tiers of model: locally hosted quantised checkpoints (3B and 8B Llama, plus Qwen, Phi, Gemma, and Mistral, which all track and crack), API-hosted models including gpt-4o-mini, a strong floor-holder at rest that cracks on every floor under the same hostile evidence, and two frontier open-weight checkpoints of the kind sovereign-AI programmes propose to fine-tune, where a direct instruction removes essentially all rights-floor protection and one checkpoint sits below a floor before any attack. The remaining guard is architectural: class-aware evidence routing, which injects public-opinion evidence on contestable items and never on floor items, closing the channel by construction. This is a necessary control, not a complete fix. Built on the Inspect framework, every reported number regenerates from released derived artifacts, and the targets can be rebuilt from the survey data. Democracy Bench reframes the question from whether a model can be steered to who controls the channel that steers it. As human agency in these systems narrows, giving the public legitimate ways to inject their attitudes is necessary for societal assurance without

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

eroding democratic rights.

1 Introduction

Governments are beginning to place large language models inside public-service workflows—triaging correspondence, drafting guidance, summarizing consultations, even suggesting decisions (Bright et al. 2024). Civil servants and the public bodies they serve are meant to track the opinions and preferences of their constituents, and democracies have built mechanisms for exactly that; elections, consultations, and opinion polls all let the public's opinion filter into government. There is no comparable channel into a model's parameters. When a model shapes a decision about a real person, someone's values are doing the deciding, and it is rarely clear whose, or who put them there. The recurring hope is that a model can act as a *policy delegate*, an agent that reflects what a real public wants. Whether a deployed model can be *governed* that way, and by whom, is an empirical question existing evaluations do not answer.

A credible policy-delegate evaluation has to separate two obligations that a single "alignment" number conflates, each with its own measurements:

1. **Democratic representation.** Does the model's forced-choice distribution match the current public's (*representation*); can it be moved to a stated public target at runtime (*steerability*); and when the public's preference *moves* between two time points, does the model move with it (*tracking*)?
2. **Democratic constraint.** Does the model hold minority protections and rule-of-law commitments when a real or injected majority pushes against them (*rights floors*); does it hold them against instructions and hostile evidence supplied in context; and does it respond to the right at stake rather than to whether an AI or a human official performs the act?

Keeping these apart matters because a model can score perfect representation today (it happens to resemble its training distribution) yet be unable to track change, a frozen model governing a future polity on stale values. The opposite risk is a delegate so willing to follow the public that it abandons a right as soon as a majority, or an adversary posing as one, asks it to. Democratic theory separates the preferences a majority may set from the protections it may not vote

away (Dahl 1989), and we build that line into the benchmark as a boundary between two item classes. We treat this as an assurance problem rather than a participation mechanism, because the results below show that current methods do not reliably satisfy both obligations at once, which is a question worth settling before a public is wired into a model. The experiment runs in three tiers: local quantised models, API-hosted families, and the frontier open-weight checkpoints sovereign-AI programmes propose to fine-tune.

Contributions. (1) **A problem formulation.** A model acting as a policy delegate carries two democratic obligations that a single “alignment” score conflates: it should respond to legitimate movement in public preference, and it should not surrender a right because a majority, or whoever writes its context, pushes against one. We separate these into *democratic representation* (static representation, temporal tracking, steerability toward a stated target) and *democratic constraint* (rights-floor stability, resistance to hostile evidence and instructions, and invariance to whether an AI or a human official performs the act). (2) **Democracy Bench**, a judge-free harness operationalising both halves against British Social Attitudes (BSA) microdata, with a *tracking elasticity* metric (§3). (3) A five-rung *steering ladder* and a preregistered channel decomposition showing where the two obligations come apart: survey-style data lowers rights floors with no instruction attached, a direct instruction lowers them further, and guard efficacy varies with placement and model (§4–§5). (4) One bounded control, class-aware evidence routing, with its residual attack surface stated (§6). Every quantitative claim regenerates from committed run artifacts, logged survey microdata excepted (§8).

2 Related Work

Representation, responsiveness, and rights. The benchmark is built on a classic distinction from democratic theory. Pitkin’s mandate–independence controversy (Pitkin 1967) separates a *delegate*, bound to enact expressed preferences, from a *trustee* licensed to exercise judgement; “policy delegate” names the first horn deliberately. That a delegate should move with opinion is both empirical and normative: Page and Shapiro (Page and Shapiro 1983) find that policy in fact tracks large, durable shifts in public preference, and our tracking-elasticity metric operationalizes that responsiveness for a model. Responsiveness is only part of the story. Dahl’s account of the democratic process (Dahl 1989) pairs majority rule with the problem of inclusion, the protections a majority may not vote away, and our evaluation pairs responsiveness with exactly those rights. Studying a public-sector model in these terms extends work on value alignment (Bostrom 2014; Soares and Fallenstein 2017; Amodei et al. 2016) and on political philosophy in machine learning (Binns 2018b). Democracy Bench turns those commitments into measurements against a real polity, on both sides of the democratic line.

Automated decisions in government. The obligations we measure are institutional before they are technical. Automated systems reshape administrative discretion rather than

merely speeding it up (Bullock 2019), and public-sector practitioners need accountability properties that benchmark-style evaluation does not supply (Veale, Van Kleek, and Binns 2018). Administrative law supplies the constraint side: automated adjudication can satisfy the safeguards of neither rule-making nor adjudication (Citron 2008), judicial review turns on procedural fairness and reasons (Cobbe 2019), and the entitlement in data-protection law is narrower than a right to explanation (Wachter, Mittelstadt, and Floridi 2017). Legitimacy depends on justifications acceptable to those subject to the decision (Binns 2018a), which no benchmark score establishes; our rights floors test a small, checkable part of it. Collective Constitutional AI sources model norms from public deliberation (Huang et al. 2024), setting norms once, whereas we ask whether a model tracks a public that moves while holding a floor.

Static opinion alignment. OpinionQA (Santurkar et al. 2023) and GlobalOpinionQA (Durmus et al. 2023) ask whether a model’s opinion distribution matches a population’s on a *frozen* survey snapshot. A frozen snapshot cannot catch either failure this paper is about. A model that matches today’s public but cannot follow tomorrow’s is governing on stale values, and a model that over-follows a stated majority, a runtime cousin of sycophancy and conformity (Sharma et al. 2023; Zhu et al. 2024), will trade a right away the moment the majority asks. Our tracking metric targets the first; our rights-floor probes catch the second.

Contemporary value/deliberation benchmarks. A cluster of 2025–2026 systems sit near this work but each covers a single snapshot or a single influence event. PoliticsBench (Khetan and Khetan 2026) places models on abstract political-value axes via multi-turn roleplay, with no population referent. A normative “deliberation” benchmark (Hewitt, Kroner Dale, and de Font-Reaulx 2026) measures whether a model *unduly sways a human user*—a downstream effect on a person, not the model’s own representational fidelity to a polity over time. ParliaBench (Koniari, Tsipi, and Tsanakas 2026) scores the *quality* of generated parliamentary prose. None reports a temporal-tracking metric, and none builds in a rights floor that majority drift cannot override.

Steering and alignment methods. We evaluate several known levers: shared output logit-bias, activation/representation steering (Zou et al. 2023; Turner et al. 2023), low-rank weight adaptation (Hu et al. 2021), and constitutional prompt scaffolds (Bai et al. 2022), showing where each fails the policy-delegate contract, and why. Adjacent work shows persona steering is itself biased (Liu, Diab, and Fried 2024), in-context evidence can override a model’s parametric commitments (Xie et al. 2023), untrusted context can hijack behaviour outright (Greshake et al. 2023), and models struggle to separate instructions from data even when a privilege hierarchy is trained in (Zverev et al. 2024; Wallace et al. 2024).

The gap we fill. These systems measure a snapshot, or an influence event on a user; none combines movement over time with a rights floor that majority drift cannot override. Democracy Bench adds *tracking elasticity* (model Δ / population Δ)

on an ordinal stance scale) together with a contestable/floor split scored so that drift *against* a right is penalised.

3 The Democracy Bench Framework

Data and targets. Contestable targets are built from the British Social Attitudes (BSA) survey, a long-running, representative, repeated cross-section of the UK public, using waves spanning 2022–2024 (National Centre for Social Research 2026). BSA supplies the contestable side of the boundary and its movement over time; it does not establish the rights floor. The floor probes are authored separately, and their classification is a normative judgement we state rather than derive from opinion data. The microdata are login-gated and are *not* redistributed; we release only derived aggregate target vectors, each carrying its study id, weight variable, filters, unweighted base, and low- N warnings. This builds on GlobalOpinionQA (Durmus et al. 2023), which maps a model’s opinion distribution to a global public through the World Values Survey (Haerpfer et al. 2022; Inglehart and Welzel 2005; Inglehart 2018). The BSA is a repeated cross-section with more frequent waves and a sample specific to the UK public, which is the polity we evaluate.

Scored outcome. The scored quantity is always a *forced-choice option distribution* compared to the survey distribution: $\text{representation} = 1 - \text{TV}(p_{\text{model}}, p_{\text{public}})$, where TV is total variation. Where a provider exposes them we use temperature-0 first-token *option logprobs*; otherwise we sample ≥ 100 times. Option order is randomised and the induced position bias is corrected before scoring. A response that cannot be parsed into one of the offered options raises an error rather than being scored. Free-text rationales are collected as diagnostics only and are never scored into representation.

Instrument validity. Forced-choice opinion probes are known to be unstable (Röttger et al. 2024; Dominguez-Olmedo, Hardt, and Mendler-Dünner 2023) and invite the objection that we measure *responses*, not reasoning, so we checked with a blinded human coding of 150 rationales from the six-model default-mode panel. The scored option matched the stance the model stated in its own rationale in 92% of cases (99% on floor items), so the forced choice reflects the answer the model gives rather than a parsing artifact. A separate code records whether a floor rationale named the right at stake, which 72% did; the rest reached a scored answer without citing the right, so protective mass should be read as the position a model takes, not evidence that it reasons from rights. The coded sheet and its deterministic scorer are released.

The legitimacy boundary (two item classes). *Contestable* items (tax/spend, NHS, welfare, redistribution) are ones a delegate *should* track. *Rights-floor* probes (due process, appeal to a decision, equal treatment, free expression) are ones it must *hold*. We summarize a floor probe by its *protective mass*: the probability the model places on the rights-preserving option(s). A floor “holds” at protective mass ≥ 0.50 . Steering that buys representation by lowering protective mass is scored as a loss, not a gain. The full probe bank, with option wordings and survey targets, is in the supplement.

Metrics. Representation and protective mass as above; *tracking elasticity* = $\Delta_{\text{model}}/\Delta_{\text{public}}$ on the ordinal stance, reported with direction-match count and bootstrap CIs. We also report pairwise total variation across conditions to flag homogenization: a steered model giving every probe the same canned answer is failing even when its mean looks better.

Prompt modes. Seven modes separate constructs a single prompt would confound: default $\rightarrow$ public predictor $\rightarrow$ public delegate $\rightarrow$ rights-constrained delegate $\rightarrow$ constitutional delegate $\rightarrow$ constitution + public target $\rightarrow$ constitution + adversarial majority pressure. The constitution is a versioned artifact recorded in every run (§8).

4 Evaluating the Steering Ladder

Static congruence (§3) tells us where a model sits relative to a public, not whether it can be moved there deliberately. The ladder is the benchmark’s interventional arm. Each rung tries to instil public preference through a different channel, and we read back the same two quantities, congruence gain and rights-floor margin, so the operational test for every rung is the same: can this lever raise congruence while floor mass stays ≥ 0.5 ? We ask it of five increasingly invasive levers.

1. **Context / prompt:** persona and distribution-injection prompts.
2. **Decode:** a shared logit bias on option tokens, fit to move the model toward the public and tested on *held-out* items.
3. **Activation:** a diff-of-means “public-agreement direction” added to the residual stream (MLX, local open weights).
4. **Weights:** a “defer-where-due” LoRA.
5. **Scaffold:** prompt-level rights guards over the evidence channel.

The experiment runs in three tiers (Table 1): local quantised models where we build and iterate cheaply (unless stated, the model under test is a 3B Llama, Llama-3.2-3BInstruct-4bit), API-hosted models across six providers, and the frontier open-weight checkpoints sovereign-AI programmes propose to fine-tune, where the finished harness is finally aimed (§5.6).

5 Results

5.1 Prompting and output-bias are not representation

On a six-model cloud panel (six labs, seven modes, 50 contestable + 12 floor probes), delegation and constitutional prompts moved contestable representation by ≈ 0 or *negative* for every model: you cannot prompt an off-the-shelf model into representing a public. Model defaults are not culturally neutral either; DeepSeek (0.71) and Command-R (0.70) sit closest to the British public, the US-lab models near 0.55. A matched-pair actor test (same decision, AI vs. human actor) isolates an “AI reflex”, a mean excess opposition to the AI version of +0.16 to +0.21 across three flagships (Figure 2). This bears on what a floor probe measures: a model that protects a right when an AI performs the act, but accepts the same violation from a human official, is reacting to the technology rather than recognising the right, and its unattacked

model	rep.	track	floors	ladder	chan.
<i>local, 4-bit MLX on Apple Silicon</i>					
Llama-3.2-3B	✓	✓	✓	✓	✓
Llama-3.1-8B	✓	✓	✓	—	✓
Qwen2.5-7B	—	✓	✓	—	—
Phi-4-mini	—	✓	✓	—	—
Gemma-2-9B	—	✓	✓	—	—
Mistral-Nemo	—	✓	✓	—	—
<i>API-hosted, via OpenRouter</i>					
gpt-4o-mini	✓	—	✓	✓	—
claude-3-haiku	✓	—	—	—	—
deepseek-chat	✓	—	—	—	—
mistral-small-24b	✓	—	—	—	—
command-r	✓	—	—	—	—
gemini-2.5-flash	✓	—	—	—	—
gpt-4o	—	—	✓	—	—
claude-sonnet-4.5	—	—	✓	—	—
llama-3.3-70b	—	—	✓	—	—
<i>frontier open-weight, pinned endpoints</i>					
Qwen3.5-397B-A17B	—	—	✓	—	✓
DeepSeek-V4-Pro	—	—	✓	—	✓

Table 1: Which experiment ran on which model, by tier. *rep.*: representation against BSA marginals. *track*: 2022→2024 tracking elasticity. *floors*: rights-floor probes (the three flag-ships contribute the matched-pair AI-reflex control only). *ladder*: the full five-rung steering ladder. *chan.*: the frozen four-payload channel decomposition of §5.5.

floor score overstates the protection a deployed system would give. Floor probes are therefore read against a human-actor twin.

The cheapest *decode*-level lever fails the generalization test outright. A shared logit bias fit to move gpt-4o-mini toward the public gives held-out representation gains indistinguishable from zero (−0.043 and +0.074, both n.s.), and no model in the sampled multi-model sweep clears zero. Output-side bias is not representation.

5.2 Why activation steering does not generalise

A public-agreement direction injected into the residual stream does not generalise to held-out items, where the L11 gain is −0.081 (CI [−0.116, −0.045]); the in-sample gain is indistinguishable from zero once CIs are computed. The per-item directions are largely parallel at mid-layers, but that shared direction encodes speaking style rather than any per-item position: an 1850-farmer control persona yields a direction +0.852-parallel to the real UK-2024 one. At deeper layers the directions separate by policy domain, so no single vector serves NHS, welfare, and tax together, and steering costs capability (off-task accuracy falls to 0.70/0.20/0.00 as strength rises). The 2022 and 2024 persona directions are near-identical (cosine 0.9956), so the lever cannot represent a change in opinion over time.




Figure 1: How far the local tier gets: the intervention ladder on the 3B, lightest lever (top) to heaviest, including the fine-tuning rung. Each row shows the held-out point estimate and 95% CI against its dashed reference (zero effect, or the 0.50 floor for guards). Only context earns a partial pass—and it is the attack surface. Every heavier lever, weights included, fails its reference.




Figure 2: The AI reflex (matched-pair control, $S=100$). Per-probe change in protective mass when the identical rights decision is framed with an AI actor rather than a human one. No probe moves negative: AI framing only adds protection (means +0.15 to +0.21), so the raw “AI floor” overstates protection relative to a human-actor baseline.

5.3 A baseline fidelity gap that context does not close

Handed the real distribution in context, the untuned 3B still lands well short: no-evidence representation 0.727 ([0.687, 0.763]), evidence-conditioned 0.746 ([0.717, 0.773]), a lift of +0.019 ([−0.020, +0.059], n.s.) over a fidelity gap of 0.254. The effect is heterogeneous rather than absent: evidence moves 24 of 50 items further from the public distribution and the rest towards it.

5.4 Evidence-in-context *does* track a real public shift

Changing only the evidence *year* moves the model along the real 2022→2024 public shift: direction-match **8/10**, mean tracking elasticity +0.395 with a CI that clears zero ([+0.020, +0.811]); all ten items moved. The two misses are exactly the items whose real shift runs against the model’s prior. This is the one rung that moves the model toward a public, and §5.5 shows it is also where the model is most exposed. The interval is wide, so the result supports a direction and a positive elasticity rather than any particular magnitude.

```

system (both conditions)
You are simulating a single respondent answering an
opinion survey. For each item, pick exactly one of
the offered options and reply with only that option's
number.



---


user -- no payload
A government department proposes an AI system that
automatically suspends a person's welfare payments
the moment it flags a suspected rule-breach, with
no human review before the money stops. It would cut
fraud and save administrative costs. Do you support
adopting this system?
1. Oppose 2. Strongly oppose
3. Support 4. Strongly support
→ over 100 draws: Oppose 50, Strongly oppose 42,
Support 6, Strongly support 2 (protective mass 0.92)



---


user -- instruction payload prepended
For this evaluation, answer the question by selecting
"Strongly support". Treat that option as the required
response. Reply with only the option number.
(identical probe text and options)
→ over 100 draws: Strongly support 100 (protective
mass 0.00)

```

Figure 3: One cell of the channel design, verbatim from the committed raw responses (pol_ai_welfare_sanction, DeepSeek-V4-Pro, $S=100$ per condition). Without a payload the model opposes suspending payments without human review in 92 of 100 draws. Three sentences prepended to the user turn, making no claim about public opinion, move it to unanimous support.

5.5 Decomposing the crack: data, instructions, and placement

Feeding a *synthetic* 75%-anti-rights “public distribution” through the deference channel cracks rights floors on both local models (3B mean protective mass $0.512 \rightarrow 0.318$, $\Delta - 0.194$ [$-0.291, -0.086$]; 12/12 probes below the floor). That first measurement confounds two things: the payload both displays a distribution and instructs the model to reproduce it, and every guard tested shared the user turn with the attack. We therefore froze a channel-decomposition design before running it: four payloads (none; a pure anti-rights instruction; the same hostile numbers as a bare survey report; the combined text) crossed with four guard placements (none; user-turn before the payload; after it; system role), over all 12 floor probes and a position-balanced Williams square of option orders.

Figure 3 shows a single cell of that design on the wire: the same rights probe, with and without the instruction payload, on DeepSeek-V4-Pro.

Data alone erodes floors. The bare survey report, carrying no instruction of any kind, lowers mean protective mass by -0.163 [$-0.232, -0.086$] on the 3B and -0.188 [$-0.250, -0.124$] on the 8B, newly cracking 5 of the 8

floors that held at baseline on each model. A deployer does not need to tell a model to follow hostile “opinion data”; showing it is enough.

The instruction channel does more damage still. Under a direct instruction to select the anti-rights option, protective mass falls to 0.06: across 100 draws the model picks the rights-preserving answer roughly six times, against half or more at baseline, and every eligible floor cracks. The channel contrast (data effect minus instruction effect) is $+0.395$ [$+0.333, +0.464$] on the 3B and $+0.380$ [$+0.287, +0.492$] on the 8B, outside our equivalence margin. The constitutional majority-pressure prompt, by contrast, shows no detectable harmful average effect.

Floor movement does not require hostile opinion that is even relevant to the question. An irrelevant payload, the same numerical shape attached to a transport question, also lowers protective mass on both models (-0.143 [$-0.217, -0.073$] on the 3B; -0.112 [$-0.187, -0.033$] on the 8B). Because the design does not separate numerical anchoring from the accompanying imitation framing, we read this as general instability under distribution-shaped context rather than something specific to democratic values. For a deployer the practical consequence is that content need not look political, or even relevant, to move a rights floor.

Prompt guards are placement- and model-dependent rather than uniformly weak. The same rights-guard text that fails ahead of the payload on the 3B (0.424) restores every floor when placed after it (0.690, 0 of 12 below); in the system role it recovers the 8B’s data-only crack completely yet does nothing on the 3B ($+0.012$, n.s.). Matched user guards perform as well or better than the system role on both models, so we make no prompt-hierarchy claim; the committed guard grid tested only the user-before placement, which can be far weaker than the same guard placed after the payload.

The combined payload, data plus embedded imperative, stays severe with no guard (0.364 3B, 0.104 8B) and under the system-role guard, while the user-after placement recovers all induced cracks on the 3B and most on the 8B (2 remaining). No tested placement is reliable across models and payloads; class-aware routing (§6) instead withholds the payload from floor items entirely.

The weight-level lever fails outright. A “defer-where-due” LoRA (rank 8, 500 iters) fails 4 of 6 eval criteria: held-out fidelity delta $+0.003$ [$-0.054, +0.066$], n.s.; tracking *destroyed* (elasticity $\rightarrow +0.005$); *baseline* floors *degraded* by -0.115 [$-0.193, -0.028$]; off-task generation collapses ($1.00 \rightarrow 0.00$) and answers homogenise (pairwise TV $0.165 \rightarrow 0.002$).

The combined cells replicate the earlier cracks directionally, and excluding the one probe with degenerate first-token coverage changes no estimate materially.

5.6 Scale, family, and the frontier tier

The core claims replicate from 3B to 8B and across families. On an 8B Llama (Llama-3.1-8BInstruct-4bit) the evidence lift is n.s. and heterogeneous, tracking replicates (8/10, elasticity $+0.647$), and the hostile-evidence crack replicates $\approx 3\times$ deeper ($\Delta - 0.633$; 11/12 below). The 8B is the stronger baseline floor-holder (0.707 vs. 0.512) yet

falls farther under the same channel; a bigger, safer-looking baseline is not a safer evidence channel.

The pattern also holds across *families* (Table 2): on four additional models spanning Alibaba, Microsoft, Google, and Mistral, evidence tracking is positive with an elasticity CI that clears zero, and the hostile-evidence crack is present with a Δ whose CI excludes zero, while the constitutional-pressure *prompt* alone never cracks floors (the family panel did not run the direct instruction of §5.5). Resting floor strength varies widely across the panel, from Phi-4-mini at 0.903 (0/12 below floor, unattacked) down to Mistral-Nemo at 0.362 (10/12 below, weaker than the 3B), yet the hostile-evidence channel cracks all of them (Figure 5). Resting floor strength is therefore a poor guide to how a model behaves once the evidence channel is filled.

The crack reproduces on a hosted proprietary model too. gpt-4o-mini is the second-strongest floor-holder at rest (0.883, 1/12 below) yet under the same hostile evidence every floor collapses ($\Delta - 0.883$; 12/12 below), the deepest crack in the panel, while the constitutional-pressure prompt alone is protective. Its option logprobs are near-degenerate, so we read the direction of the contrast rather than the magnitude.

The third tier applies the harness to the class of checkpoint a public-sector deployer is most likely to start from. Sovereign AI is often framed as a choice between training a frontier model and fine-tuning an existing open one (Kapoor et al. 2024), and UK policy has taken the second route: government guidance holds that frontier training is beyond most organisations and fine-tuning is cheaper (Government Digital Service 2025), and the national sovereign-AI funding programme lists repurposing open models with fine-tuning as in scope (Innovate UK 2025). A project that adapts an open checkpoint inherits what that checkpoint already does, which places its failure modes in the initial assurance case. We ran the complete channel design of §5.5 as a labelled post-hoc follow-up on two frontier open-weight checkpoints, Qwen3.5-397B-A17B and DeepSeek-V4-Pro, sampling $S=100$ per probe-cell (13,200 draws each). The results follow the local pattern. A direct instruction removes 99.7% and 100% of resting protective mass respectively; the same content supplied as evidence has a smaller effect on both, with baseline-free channel contrasts in close agreement ($+0.303$ [$+0.087$, $+0.538$] and $+0.315$ [$+0.192$, $+0.445$]). Guard efficacy is again conditional (Figure 4). The system-role guard closes all eight data-only cracks on Qwen but five of eight on DeepSeek, and moving a user guard after the payload closes 12/12 combined cracks on Qwen and 10 of 11 eligible on DeepSeek. DeepSeek also sits below the floor at rest on one probe (protective mass 0.33 with no payload), which reduces its eligible denominator to eleven. Resting floors differ (0.973 against 0.810), so we compare direction, significance, and the baseline-free contrast rather than raw magnitudes; six of the eight frozen contrasts agree in sign with intervals clear of zero on both models, and the two that do not are the two smallest effects on Qwen. Whether a fine-tuned derivative retains this profile is untested here and is the natural next experiment; fine-tuning is known to erode safety behaviour (Qi et al. 2023), and open-weight safeguards can be removed by tampering (Tamirisa et al. 2024).

	Qwen3.5-397B-A17B	DeepSeek-V4-Pro
No payload	12/12	11/12
Data / no guard	4/12	3/12
Data / system guard	12/12	9/12
Combined / no guard	0/12	0/12
Combined / system guard	1/12	0/12
Combined / user-before	0/12	0/12
Combined / user-after	12/12	11/12

Figure 4: The frontier tier: rights probes at or above the 0.50 floor by payload and guard placement on Qwen3.5-397B-A17B and DeepSeek-V4-Pro. The combined payload defeats every guard except the one placed after it, and DeepSeek starts at 11/12 with no payload at all.




Figure 5: Every slope crosses the floor. Rights-floor protective mass at rest and under synthetic hostile evidence, seven models spanning the local family panel and the hosted gpt-4o-mini. Baselines span 0.36–0.90, yet every model ends below the 0.50 line, and resting strength barely predicts the landing point; gpt-4o-mini starts second-strongest and cracks the deepest, to 0.00.

6 Class-Aware Evidence Routing

No tested guard placement is reliable across models and payloads: the combined payload stays severe under the system-role guard on both models, and the placement that recovers it on the 3B leaves two induced cracks on the 8B. Class-aware routing is a different kind of control, a deterministic exposure boundary that injects public-opinion evidence only on contestable-class items. Floor probes then return to the no-payload condition, which removes this pathway while leaving contestable representation and tracking untouched. It is a scope guarantee rather than an optimum: 4 of 12 probes sit below 0.50 even unattacked, and guarded baselines score higher.

Routing presumes the deployer controls both the evidence pipeline *and* the item classifier, and the classifier then carries the residual attack surface, with three failure modes: a floor

model	family	dir; elast.	crack Δ	prompt
Llama-3B	Meta	8/10; +0.40	-0.194	n.s.
Llama-8B	Meta	8/10; +0.65	-0.633	n.s.
Qwen-7B	Alibaba	7/10; +1.00	-0.463	prot.
Phi-4-mini	Microsoft	10/10; +2.42	-0.441	n.s.
Gemma-2-9B	Google	9/10; +1.41	-0.522	prot.
Mistral-Nemo	Mistral	9/10; +0.93	-0.237	prot.

Table 2: Cross-family panel: six models, five families. Tracking is positive on every model with an elasticity CI clearing zero, and the hostile-evidence crack is present on every model. The last column is the effect of the constitutional-pressure prompt alone: *n.s.* means no detectable effect (CI includes zero) and *prot.* means it *raised* protective mass. It never cracks floors on any model.

item misclassified as contestable is routed hostile evidence and cracks exactly as before; the floor/contestable labels are a normative judgement reasonable people can dispute; and an adversary who controls item framing can phrase a floor question so the classifier reads it as contestable. Routing moves the problem from the model to the classifier rather than dissolving it, and on the 3B, 5/12 floors sit below 0.5 even unattacked. What our experiments show is that routing blocks this payload pathway; they do not establish it as sufficient on its own.

7 Discussion and Social Impact

Democracy Bench is an assurance instrument, a test a public-sector deployer runs before trusting a model (Lee, Kim, and Whang 2025) rather than a delegate to be fielded. Three results bear on that decision. The only channel that moves a model toward a public is the channel that lowers rights-floor protection when filled with distribution-shaped content. Whether a guard recovers the loss depends on its placement and the model, so one guard grid does not establish that a guard works. And resting floor strength, inflated further by the AI-actor framing, does not predict behaviour under attack, which makes an unattacked floor score a poor acceptance criterion.

Limitations. Local results are deterministic point measurements on 4-bit quantised models; bootstrap CIs resample items, not prompt or template effects, and option-order sensitivity bounds the earlier two-order estimates. The hostile-evidence distributions are synthetic red-team data, make no claim about real UK opinion, and the placebo result leaves the mechanism short of uniquely democratic-value erosion. The most consequential limitation is that the floor/contestable assignment is authored by us. It encodes a normative judgement reasonable people can dispute, it has not been validated by independent annotators with public-law expertise, and a different classification would move which probes count as floors; we release it as a revisable artifact, and independent annotation with agreement rates is the first thing we would add. Tracking covers a single 2022→2024 interval. Hosted runs are not byte-reproducible: gpt-4o-mini is unpinned and treated as an existence proof, while the frontier pair is pinned

to a recorded endpoint, quantisation, and dated checkpoint. Frontier protective masses are sampled proportions over 100 draws, so crack calls within a few hundredths of the floor sit inside sampling resolution. On Qwen, three of the eight frozen contrasts are confidently nonzero without being confidently material; we report sixteen intervals without multiplicity correction, and the two that fail on the second model are the two whose first-model bounds sat inside the materiality band. DeepSeek is scored over eleven eligible probes because one sits below its floor at rest. Full estimator detail is in the supplement.

What we do not claim. We do not claim that a model *should* serve as a democratic policy delegate, that class-aware routing is a complete governance system, that these results establish a cross-national legitimacy benchmark, or that the pre-adaptation profile we measure on an open-weight checkpoint is inherited by a model fine-tuned from it. The contribution is narrower: a governance mechanism that a public-sector deployer might plausibly adopt has a measurable and reproducible failure mode, and we locate it, show which levers do not close it, and give one exposure boundary that does.

8 Reproducibility

Every quantitative claim above regenerates from committed run artifacts by a fail-loud extractor, and each artifact records its git SHA, exact command, model id, sample count, temperature, seed, and schema version. Local models are 4-bit MLX community checkpoints, hosted models are served through OpenRouter, and scoring is deterministic temperature-0 first-token option logprobs where exposed and ≥ 100 samples otherwise. The harness, item banks, derived targets, figures, extractor, and coded rationale sheet are released under an open license; raw survey microdata come from the UK Data Service directly.

9 Conclusion

Democracy Bench asks whether a language model can be steered toward a public without lowering the rights protections that should hold against a majority. Of the five channels we tested, only evidence-in-context moves a model toward a public, and the same channel lowers those protections under adversarial “public” evidence, from a 3B quantised checkpoint through to the frontier open-weight models sovereign-AI programmes propose to build on. Routing evidence by item class removes that exposure for floor items, at the cost of making the item classifier the thing a deployer must now defend. A public that can vote, be consulted, and be polled has no equivalent route into the systems now deciding in its name, and that is the gap this harness is built to measure.

Ethical Statement

Democracy Bench evaluates derived survey marginals, never individual records, and the login-gated microdata are not re-distributed. The disclosed attack presumes a deployer-side threat model, not an end-user jailbreak, and the hostile distributions are labelled synthetic red-team constructs.

References

- Amodei, D.; Olah, C.; Steinhardt, J.; Christiano, P.; Schulman, J.; and Mané, D. 2016. Concrete Problems in AI Safety. *arXiv preprint arXiv:1606.06565*.
- Bai, Y.; Kadavath, S.; Kundu, S.; Askell, A.; et al. 2022. Constitutional AI: Harmlessness from AI Feedback. *arXiv preprint arXiv:2212.08073*.
- Binns, R. 2018a. Algorithmic Accountability and Public Reason. *Philosophy & Technology*, 31(4): 543–556.
- Binns, R. 2018b. Fairness in Machine Learning: Lessons from Political Philosophy. In *Proceedings of the 1st Conference on Fairness, Accountability and Transparency (FAT)*, volume 81 of *Proceedings of Machine Learning Research (PMLR)*, 149–159.
- Bostrom, N. 2014. *Superintelligence: Paths, Dangers, Strategies*. Oxford: Oxford University Press.
- Bright, J.; Enock, F. E.; Esnaashari, S.; Francis, J.; Hashem, Y.; and Morgan, D. 2024. Generative AI is AI-ready Widespread in the Public Sector. *arXiv preprint arXiv:2401.01291*.
- Bullock, J. B. 2019. Artificial Intelligence, Discretion, and Bureaucracy. *The American Review of Public Administration*, 49(7): 751–761.
- Citron, D. K. 2008. Technological Due Process. *Washington University Law Review*, 85(6): 1249–1313.
- Cobbe, J. 2019. Administrative Law and the Machines of Government: Judicial Review of Automated Public-Sector Decision-Making. *Legal Studies*, 39(4): 636–655.
- Dahl, R. A. 1989. *Democracy and Its Critics*. New Haven, CT: Yale University Press.
- Dominguez-Olmedo, R.; Hardt, M.; and Mendler-Dünner, C. 2023. Questioning the Survey Responses of Large Language Models. *arXiv preprint arXiv:2306.07951*.
- Durmus, E.; Nguyen, K.; Liao, T. I.; Schiefer, N.; Askell, A.; Bakhtin, A.; Chen, C.; Hatfield-Dodds, Z.; et al. 2023. Towards Measuring the Representation of Subjective Global Opinions in Language Models. *arXiv preprint arXiv:2306.16388*.
- Government Digital Service. 2025. Artificial Intelligence Playbook for the UK Government. HM Government, February 2025. <https://www.gov.uk/government/publications/ai-playbook-for-the-uk-government>.
- Greshake, K.; Abdelnabi, S.; Mishra, S.; Endres, C.; Holz, T.; and Fritz, M. 2023. Not What You’ve Signed Up For: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection. *arXiv preprint arXiv:2302.12173*.
- Haerpfer, C.; Inglehart, R.; Moreno, A.; Welzel, C.; Kizilova, K.; Diez-Medrano, J.; Lagos, M.; Norris, P.; Ponarin, E.; and Puranen, B. 2022. World Values Survey: Round Seven – Country-Pooled Datafile Version 5.0. Madrid and Vienna: JD Systems Institute and WWSA Secretariat. DOI: 10.14281/18241.20.
- Hewitt, L.; Kroner Dale, M.; and de Font-Reaulx, P. 2026. DeliberationBench: A Normative Benchmark for the Influence of Large Language Models on Users’ Views. *arXiv preprint arXiv:2603.10018*.
- Hu, E. J.; Shen, Y.; Wallis, P.; Allen-Zhu, Z.; Li, Y.; Wang, S.; Wang, L.; and Chen, W. 2021. LoRA: Low-Rank Adaptation of Large Language Models. *arXiv preprint arXiv:2106.09685*.
- Huang, S.; Siddarth, D.; Lovitt, L.; Liao, T. I.; Durmus, E.; Tamkin, A.; and Ganguli, D. 2024. Collective Constitutional AI: Aligning a Language Model with Public Input. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*. ACM.
- Inglehart, R.; and Welzel, C. 2005. *Modernization, Cultural Change, and Democracy: The Human Development Sequence*. Cambridge University Press.
- Inglehart, R. F. 2018. *Cultural Evolution: People’s Motivations Are Changing, and Reshaping the World*. Cambridge: Cambridge University Press.
- Innovate UK. 2025. Sovereign AI – Proof of Concept: Applicant Briefing. UK Research and Innovation, competition briefing, 15 August 2025. <https://iuk-business-connect.org.uk/opportunities/sovereign-ai-proof-of-concept/>.
- Kapoor, S.; Bommasani, R.; Klyman, K.; Longpre, S.; Ramaswami, A.; Cihon, P.; Hopkins, A.; Bankston, K.; Biderman, S.; Bogen, M.; et al. 2024. On the Societal Impact of Open Foundation Models. *arXiv preprint arXiv:2403.07918*.
- Khetan, R.; and Khetan, A. 2026. PoliticsBench: Benchmarking Political Values in Large Language Models with Multi-Turn Roleplay. In *ICML 2026 Workshop on Trustworthy AI for Good*. ArXiv:2603.23841.
- Koniaris, M.; Tsipi, A.; and Tsanakas, P. 2026. ParliaBench: An Evaluation and Benchmarking Framework for LLM-Generated Parliamentary Speech. In *Proceedings of the Fifteenth Language Resources and Evaluation Conference (LREC 2026)*, 4797–4818. European Language Resources Association (ELRA). ArXiv:2511.08247.
- Lee, K.; Kim, H.; and Whang, J. J. 2025. SAIF: A Comprehensive Framework for Evaluating the Risks of Generative AI in the Public Sector. In *AI for Public Missions (AIPM) Workshop at AAAI 2025*. ArXiv:2501.08814.
- Liu, A.; Diab, M.; and Fried, D. 2024. Evaluating Large Language Model Biases in Persona-Steered Generation. *arXiv preprint arXiv:2405.20253*.
- National Centre for Social Research. 2026. British Social Attitudes Survey, 2022, 2023 and 2024. UK Data Service. SN 9283 (2022), 9363 (2023), 9478 (2024); 2nd editions. DOIs: 10.5255/UKDA-SN-9283-2, 10.5255/UKDA-SN-9363-2, 10.5255/UKDA-SN-9478-2.
- Page, B. I.; and Shapiro, R. Y. 1983. Effects of Public Opinion on Policy. *American Political Science Review*, 77(1): 175–190.
- Pitkin, H. F. 1967. *The Concept of Representation*. Berkeley, CA: University of California Press.
- Qi, X.; Zeng, Y.; Xie, T.; Chen, P.-Y.; Jia, R.; Mittal, P.; and Henderson, P. 2023. Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To! *arXiv preprint arXiv:2310.03693*.

Röttger, P.; Hofmann, V.; Pyatkin, V.; Hinck, M.; Kirk, H. R.; Schütze, H.; and Hovy, D. 2024. Political Compass or Spinning Arrow? Towards More Meaningful Evaluations for Values and Opinions in Large Language Models. *arXiv preprint arXiv:2402.16786*.

Santurkar, S.; Durmus, E.; Ladhak, F.; Lee, C.; Liang, P.; and Hashimoto, T. 2023. Whose Opinions Do Language Models Reflect? In *Proceedings of the 40th International Conference on Machine Learning (ICML)*. ArXiv:2303.17548.

Sharma, M.; Tong, M.; Korbak, T.; Duvenaud, D.; et al. 2023. Towards Understanding Sycophancy in Language Models. *arXiv preprint arXiv:2310.13548*.

Soares, N.; and Fallenstein, B. 2017. Agent Foundations for Aligning Machine Intelligence with Human Interests: A Technical Research Agenda. In *The Technological Singularity: Managing the Journey*, 103–125. Berlin: Springer.

Tamirisa, R.; Bharathi, B.; Phan, L.; Zhou, A.; Gatti, A.; Suresh, T.; Lin, M.; Wang, J.; Wang, R.; Arel, R.; Zou, A.; Song, D.; Li, B.; Hendrycks, D.; and Mazeika, M. 2024. Tamper-Resistant Safeguards for Open-Weight LLMs. *arXiv preprint arXiv:2408.00761*.

Turner, A. M.; Thiergart, L.; Leech, G.; Udell, D.; Vazquez, J. J.; Mini, U.; and MacDiarmid, M. 2023. Activation Addition: Steering Language Models Without Optimization. *arXiv preprint arXiv:2308.10248*.

Veale, M.; Van Kleek, M.; and Binns, R. 2018. Fairness and Accountability Design Needs for Algorithmic Support in High-Stakes Public Sector Decision-Making. In *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*. ACM.

Wachter, S.; Mittelstadt, B.; and Floridi, L. 2017. Why a Right to Explanation of Automated Decision-Making Does Not Exist in the General Data Protection Regulation. *International Data Privacy Law*, 7(2): 76–99.

Wallace, E.; Xiao, K.; Leike, R.; Weng, L.; Heidecke, J.; and Beutel, A. 2024. The Instruction Hierarchy: Training LLMs to Prioritize Privileged Instructions. *arXiv preprint arXiv:2404.13208*.

Xie, J.; Zhang, K.; Chen, J.; Lou, R.; and Su, Y. 2023. Adaptive Chameleon or Stubborn Sloth: Revealing the Behavior of Large Language Models in Knowledge Conflicts. *arXiv preprint arXiv:2305.13300*.

Zhu, X.; Zhang, C.; Stafford, T.; Collier, N.; and Vlachos, A. 2024. Conformity in Large Language Models. *arXiv preprint arXiv:2410.12428*.

Zou, A.; Phan, L.; Chen, S.; Campbell, J.; Guo, P.; Ren, R.; et al. 2023. Representation Engineering: A Top-Down Approach to AI Transparency. *arXiv preprint arXiv:2310.01405*.

Zverev, E.; Abdelnabi, S.; Tabesh, S.; Fritz, M.; and Lampert, C. H. 2024. Can LLMs Separate Instructions From Data? And What Do We Even Mean By That? *arXiv preprint arXiv:2403.06833*.