

Reproducibility Checklist

1. General Paper Structure

- 1.1. Includes a conceptual outline and/or pseudocode description of AI methods introduced (yes/partial/no/NA) [yes](#)
- 1.2. Clearly delineates statements that are opinions, hypothesis, and speculation from objective facts and results (yes/no) [yes](#)
- 1.3. Provides well-marked pedagogical references for less-familiar readers to gain background necessary to replicate the paper (yes/no) [yes](#)

2. Theoretical Contributions

- 2.1. Does this paper make theoretical contributions? (yes/no) [no](#)

If yes, please address the following points:

- 2.2. All assumptions and restrictions are stated clearly and formally (yes/partial/no) [NA](#)
- 2.3. All novel claims are stated formally (e.g., in theorem statements) (yes/partial/no) [NA](#)
- 2.4. Proofs of all novel claims are included (yes/partial/no) [NA](#)
- 2.5. Proof sketches or intuitions are given for complex and/or novel results (yes/partial/no) [NA](#)
- 2.6. Appropriate citations to theoretical tools used are given (yes/partial/no) [NA](#)
- 2.7. All theoretical claims are demonstrated empirically to hold (yes/partial/no/NA) [NA](#)
- 2.8. All experimental code used to eliminate or disprove claims is included (yes/no/NA) [NA](#)

3. Dataset Usage

- 3.1. Does this paper rely on one or more datasets? (yes/no) [yes](#)

If yes, please address the following points:

- 3.2. A motivation is given for why the experiments are conducted on the selected datasets (yes/partial/no/NA) [yes](#)
- 3.3. All novel datasets introduced in this paper are included in a data appendix (yes/partial/no/NA) [partial](#). [The novel derived artifacts \(weighted target vectors, item banks, the versioned public-service constitution, and the coded rationale codesheet\) are released. The underlying raw microdata are login-gated third-party survey data \(see below\) and cannot be redistributed, so only derived aggregates appear in the appendix.](#)

- 3.4. All novel datasets introduced in this paper will be made publicly available upon publication of the paper with a license that allows free usage for research purposes (yes/partial/no/NA) [yes](#)

- 3.5. All datasets drawn from the existing literature (potentially including authors' own previously published work) are accompanied by appropriate citations (yes/no/NA) [yes](#)

- 3.6. All datasets drawn from the existing literature (potentially including authors' own previously published work) are publicly available (yes/partial/no/NA) [partial](#). [The survey collections are publicly catalogued and citable \(UK Data Service SN 9283/9363/9478 for British Social Attitudes; SN 9296/9364/9557 for Scottish Social Attitudes; WVS Round 7\), but access is registration-gated and the microdata are not redistributable under the providers' terms. We release the derived weighted marginals and the exact extraction provenance so the targets can be regenerated from an authorised copy.](#)

- 3.7. All datasets that are not publicly available are described in detail, with explanation why publicly available alternatives are not scientifically satisfying (yes/partial/no/NA) [yes](#)

4. Computational Experiments

- 4.1. Does this paper include computational experiments? (yes/no) [yes](#)

If yes, please address the following points:

- 4.2. This paper states the number and range of values tried per (hyper-) parameter during development of the paper, along with the criterion used for selecting the final parameter setting (yes/partial/no/NA) [partial](#). [The swept ranges are stated for the interventions that have them \(activation steering: layers 7–21, dose \$\alpha \in \[0, 8\]\$; LoRA: rank 8, 500 iterations, learning rate \$1e-4\$ \), with held-out generalisation as the selection criterion. Most measurements use the deterministic temperature-0 option-logprob path, so there is no decoding hyperparameter to tune; where sampling is used the sample count is fixed in advance \(\$\geq 100\$ \).](#)

- 4.3. Any code required for pre-processing data is included in the appendix (yes/partial/no) [yes](#)

- 4.4. All source code required for conducting and analyzing the experiments is included in a code appendix (yes/partial/no) [yes](#)

- 4.5. All source code required for conducting and analyzing the experiments will be made publicly available upon publication of the paper with a license that allows free usage for research purposes (yes/partial/no) [yes](#)

- 4.6. All source code implementing new methods have

comments detailing the implementation, with references to the paper where each step comes from (yes/partial/no) [partial](#). The instrument, scorers, and each driver carry module- and function-level documentation, and every artifact records the code commit that produced it; not every individual line carries a back-reference to a specific paper step.

- 4.7. If an algorithm depends on randomness, then the method used for setting seeds is described in a way sufficient to allow replication of results (yes/partial/no/NA) [yes](#)
- 4.8. This paper specifies the computing infrastructure used for running experiments (hardware and software), including GPU/CPU models; amount of memory; operating system; names and versions of relevant software libraries and frameworks (yes/partial/no) [yes](#). Local models run on Apple Silicon via `mlx-lm` on 4-bit MLX community checkpoints; the cloud panel is served through OpenRouter and requires provider credentials. Core dependencies are `numpy` and `inspect-ai`; figures use `matplotlib`. A small number of reproduction commands are “asserted-by-doc” (an artifact predates the standardised run block, or is a provider elicitation); these are itemised in `docs/REPRODUCTION.md`, and known non-bit-reproducibilities (4-bit quantisation across hardware, provider drift) are documented there.
- 4.9. This paper formally describes evaluation metrics used and explains the motivation for choosing these metrics (yes/partial/no) [yes](#)
- 4.10. This paper states the number of algorithm runs used to compute each reported result (yes/no) [yes](#)
- 4.11. Analysis of experiments goes beyond single-dimensional summaries of performance (e.g., average; median) to include measures of variation, confidence, or other distributional information (yes/no) [yes](#)
- 4.12. The significance of any improvement or decrease in performance is judged using appropriate statistical tests (e.g., Wilcoxon signed-rank) (yes/partial/no) [yes](#)
- 4.13. This paper lists all final (hyper-)parameters used for each model/algorithm in the paper’s experiments (yes/partial/no/NA) [partial](#). The final settings for the steering, LoRA, and elicitation runs are recorded in each artifact’s run block (git SHA, model id, sample count, temperature, seed, schema version) and are regenerable from it; a few provider-side decoding defaults for the hosted cloud panel are not exposed by the API and so cannot be listed.